

tSNE

t-distribution stochastic neighbor embedding
van der Maaten & Hinton 2008

based on Hinton & Roweis 2002 SNE

Non-Linear Dimensionality Reduction algo.

פוסט דאטא סיינס (ביאנא)
 האט חשוף אלגאריטם (למשה)
 זיכר אלט 20 יאר (20)
 ויכח בן האלגאריטם

1) "צוק" האט איהם באקאנט
 2) ער האט באקאנט איהם
 D_{KL}

Data $X: X_i \in \mathbb{R}^P$
 $\frac{SNE}{}$
 למקרה דאטא
 High-dim. data

Distances: $\forall i, j \quad d_{ij} = \|X_i - X_j\|$
 אלגאריטם דאטא
 למקרה
 d_{ij}

$$P_{j|i} = \frac{\exp(-d_{ij}^2 / 2\sigma_i^2)}{\sum_{k \neq i} \exp(-d_{ik}^2 / 2\sigma_i^2)}$$

פה, σ_i זו סכמת הניקון של הנתון (log)
 ו- P_i היא ההסתברות לזכור את i .
 (למעשה זהו מידת האי-ודאות)

איך נקבעת σ_i ? היסטוריה והנחות
 למשל, P_i של P_i כלשהו
 "לפי ההנחות האחרות"

$$\text{Perplexity} = 2^{H(P_i)}$$

כלל האגרוגים של הנתון

$$H(P_i) = - \sum_j P_{ji} \cdot \log_2 P_{ji}$$

(3-5) $\text{Perplexity} / \text{כאילו} / \text{כאילו}$ נראה
 $2^{H(P_i)}$ ו- σ_i נראה x_i של σ_i

small $\sigma_i \rightarrow$ "less neighbors" \rightarrow more predictable transitions
 lower entropy \rightarrow lower perplexity

higher $\sigma_i \rightarrow \dots \rightarrow$ higher perplexity

Embedding find low-dimensional representation
 $Y: Y_i \in \mathbb{R}^d \quad d \ll p \ll X \text{ size}$

(t-SNE or GL) SNE with points
 Gauss. Kernel, $\sigma = \frac{1}{\sqrt{2}}$ י ש המרחק הנ"ל
 $q_{j|i} = \frac{\exp(d'_{ij})}{\sum_k \exp(d'_{ik})} \quad d'_{ij} = \|y_i - y_j\|$

while minimizing the (KL) divergence of P_i, Q_i

score (for SNE): $C = \sum_i KL(P_i \| Q_i) = \sum_i \sum_j p_{j|i} \cdot \log \frac{p_{j|i}}{q_{j|i}} \quad (*)$

! Gradient Descent ? C מנסים להקטין

$$\frac{\partial C}{\partial y_i} = 2 \sum_j \underbrace{(p_{j|i} - q_{j|i})}_{\text{הפרש ההסתברויות}} \underbrace{(y_i - y_j)}_{\text{הפרש הנקודות}}$$

Init. with random points (sampled from Gauss., $\sigma = \frac{1}{\sqrt{2}}$)

Use momentum

$$Y^{(t)} = Y^{(t-1)} + \eta \underbrace{\frac{\partial C}{\partial Y}}_{\text{gradient}} + \underbrace{\alpha(t)}_{\text{momentum}} \cdot [Y^{(t-1)} - Y^{(t-2)}]$$

learning rate $\rightarrow \eta$

Comment 1 - Why KL-divergence?

Kullback-Leibler (Relative Entropy)

In Inf. theory: expected #bits (per sample)
when compressing samples from P using
code for Q (compared to a code for P)

In Bayesian Inference: Info. gained when moving from
a prior dist Q to a prior P (for data from P).

$$D_{KL}(P \parallel Q) = - \sum_i P_i \log Q_i + \sum_i P_i \log P_i = \underbrace{H(P, Q)}_{\text{cross entropy}} - \underbrace{H(P)}_{\text{entropy}}$$

Symmetric version: Jensen-Shannon Divergence

$$D_{JS}(P \parallel Q) = D_{JS}(Q \parallel P) = \frac{1}{2} D(P \parallel M) + \frac{1}{2} D(Q \parallel M)$$
$$M = \frac{P+Q}{2}$$

$Q + P \rightarrow M$ \rightarrow P and Q are both \rightarrow M

Comment 2 - So Why KL? (for SNE)

large $P_{j|i}$ are important. w/ small $q_{j|i} \rightarrow$ BIG penalty

small $P_{j|i}$, even if large $q_{j|i}$ yield small penalty

Moving to t-SNE

Solving two problems (1) Outliers
(2) Crowding

(1) X_i outlier \rightarrow add noise to C (far from all)

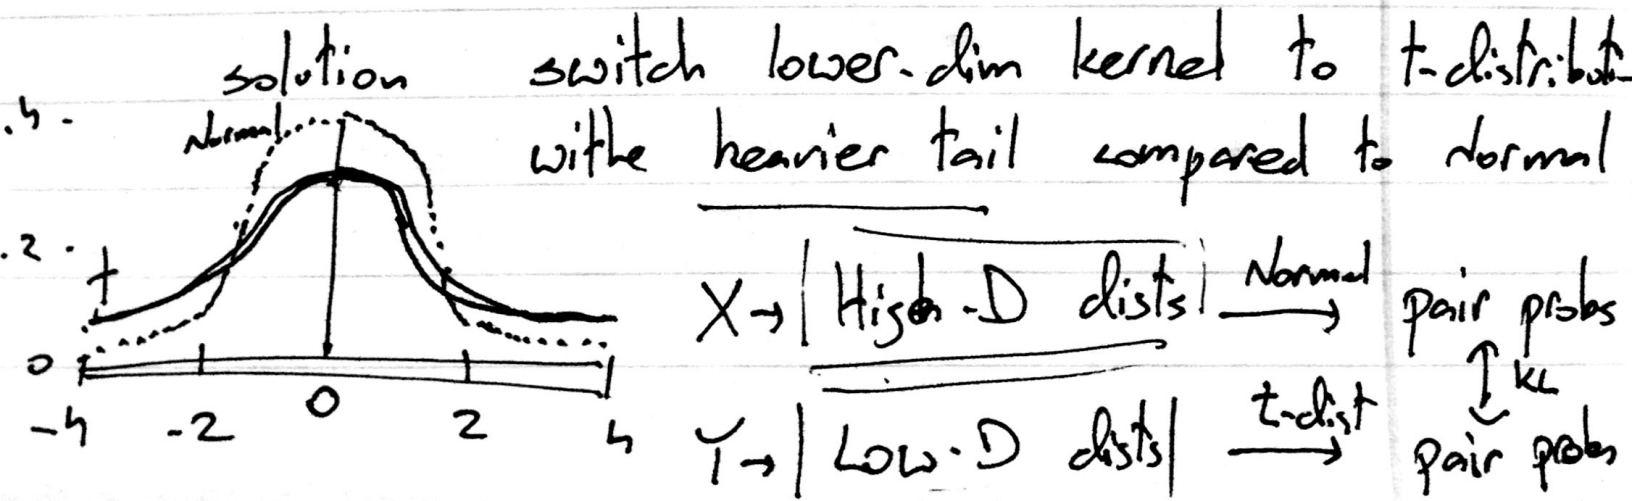

$$\text{Update: } C = KL(P \parallel Q) = \sum_{ij} p_{ij} \log p_{ij} / q_{ij}$$

Hence: move from N dist. p_{ij} to one p_{ij} (pairwise)

$$p_{ij} = \frac{p_{ij} + p_{ij}}{2 \cdot N}$$

\rightarrow large p_{ij} (= close points) get stronger

(2) High-to-low dimen. embeddings expon. reduce the "space" for neighboring points. Causing under-est. for ~~smaller~~ close points, and over-est. of dist for distant



Students t-distribution

$$\text{PDF: } f(t) = \frac{\Gamma(\frac{\nu+1}{2})}{\sqrt{\nu\pi} \Gamma(\frac{\nu}{2})} \cdot \left(1 + \frac{t^2}{\nu}\right)^{-\frac{\nu+1}{2}}, \text{ w/ param } \nu$$

$$\Gamma(x) = (x-1)!, \quad \Gamma(1) = 1, \quad \Gamma(\frac{1}{2}) = \sqrt{\pi}.$$

$$\text{w/ } \nu=1 \quad f(t) = \frac{\Gamma(1)}{\sqrt{\pi} \cdot \Gamma(\frac{1}{2})} (1+t^2)^{-1} = \frac{(1+t^2)^{-1}}{\pi}$$

in t-SNE Q is defined with a t-dist kernel

$$q_{ij} = \frac{(1 + \|y_i - y_j\|^2)^{-1}}{\sum_{k \neq l} (1 + \|y_k - y_l\|^2)^{-1}}$$

Gradient Descent $\frac{\partial C}{\partial y_i}$ $\frac{\partial}{\partial y_i} \sum_j q_{ij} (y_i - y_j) (1 + \|y_i - y_j\|^2)^{-1}$

$$\frac{\partial C}{\partial y_i} = 4 \sum_j \underbrace{(q_{ij} - q_{ji})}_{\text{HiD-LowD}} \underbrace{(y_i - y_j)}_{\text{כ"ס}} \underbrace{(1 + \|y_i - y_j\|^2)^{-1}}_{\text{היכנסת היכנסת}}$$

[for speed-up, use Barnes-Hut simulation
where neighboring points are grouped to compute forces]